

# Task 1

---

1.

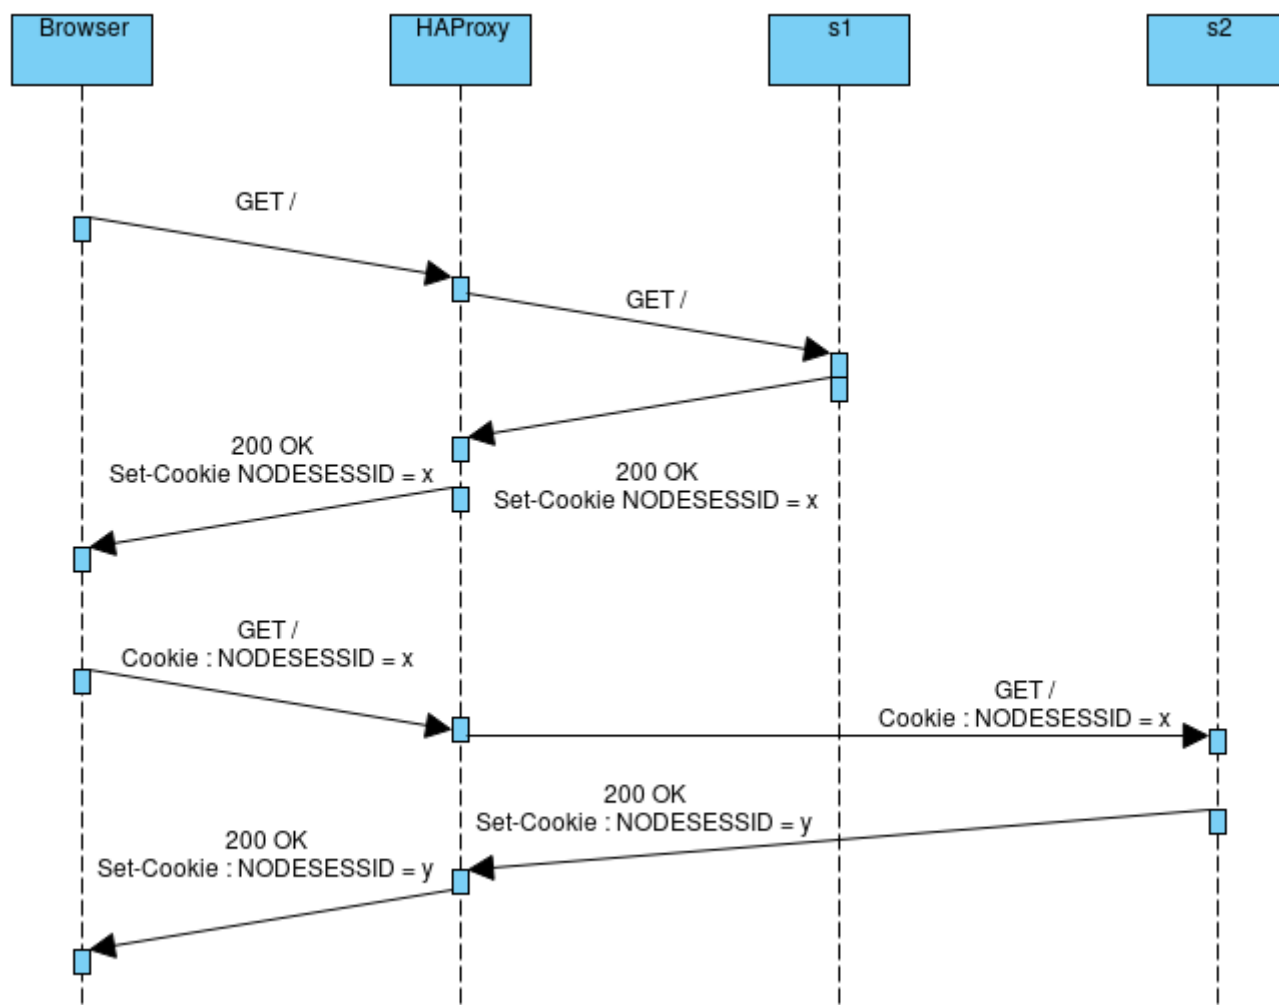
We can see that the proxy handles requests like this : It will address the requests on each server sequentially. It means that one time it will be sent to one server and the next request coming will go threw the other server. The NODESESSID is changing at each request which is not so surpirsing because each time we visit the site or actualize we have an ID which is invalid for the server requested (because the precedent response came from the other server).

2.

We expect the load balancer to address requests coming in with a session ID to the server who knows this session (so the server who created the session). If not, the session management can't be done.

3.

Here is our sequence diagram for task 1 (roundrobin). This sequence diagram shows 2 requests by the same user, we see that the behaviour is not correct because both requests don't go on the same server as expected.



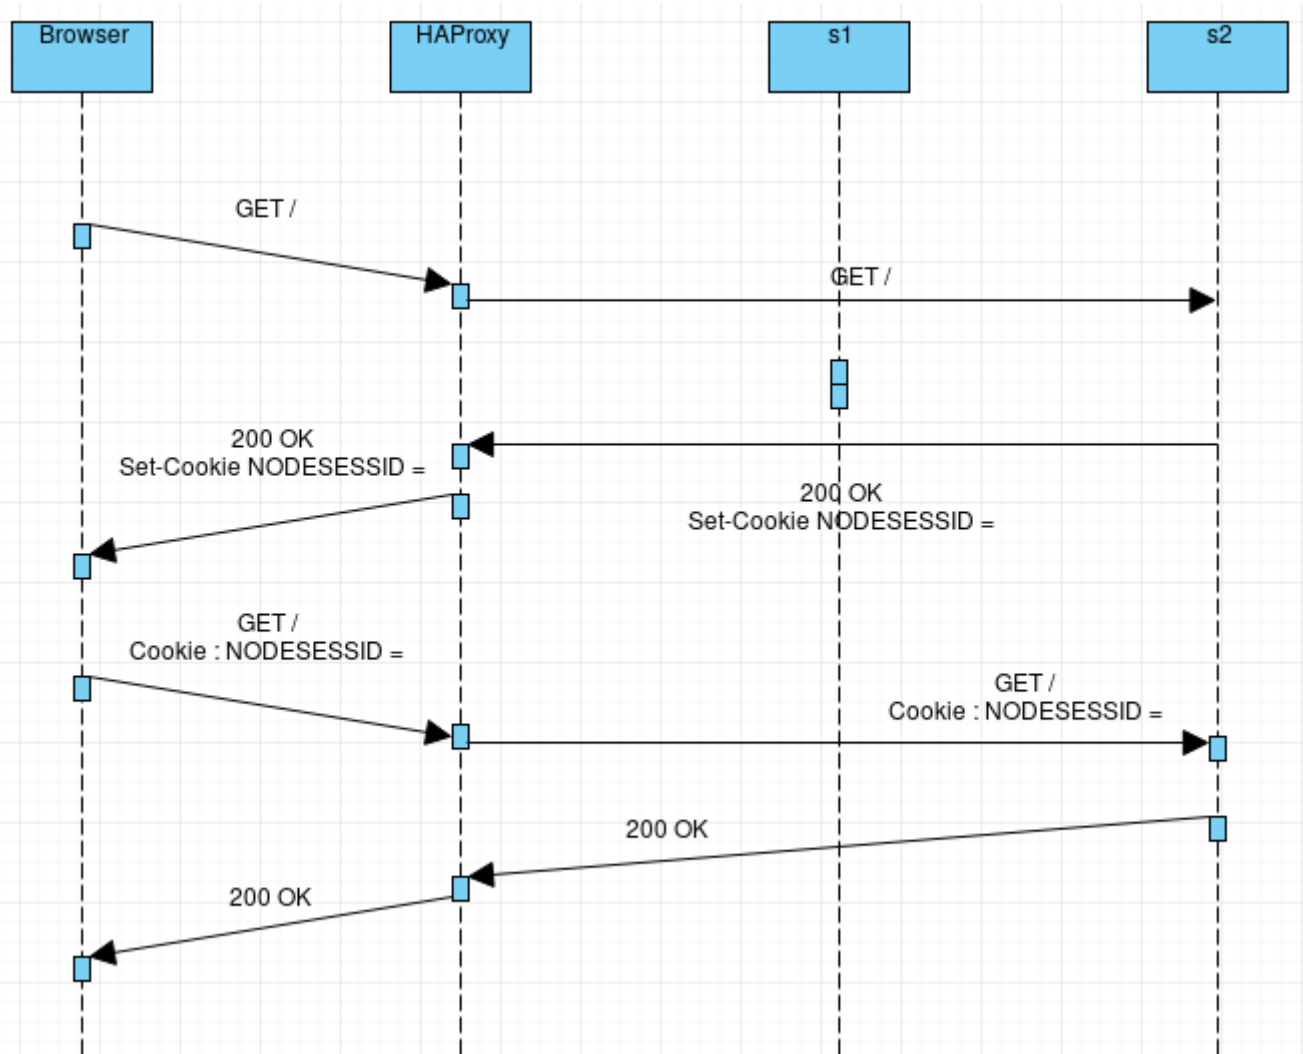
4.

Here is the test in JMeter, we can see that assertions are false because the counters doesn't increment, in fact th server don't know the users coming because of balancing strategy used here.



5.

Here we see the sequence diagram for this task, we can see that because there's only one server running it responds and knows each user (because he serves all the requests), and so the counter increments correctly.



We can see the JMeter report too :



**Rapport consolidé**

Nom : Summary Report

Commentaires :

Écrire les résultats dans un fichier ou lire les résultats depuis un fichier CSV/JTL

Nom du fichier :  Parcourir... Uniquement : Erreurs Succès Configurer

Libellé	# Echantillons	Moyenne	Min	Max	Ecart type	% Erreur	Débit	Ko/sec reus	Ko/sec émis	Moy. octets
GET /	100	6	3	22	4.05	0.00%	117.2/sec	42.61	24.96	372.1
S1 reached	100	0	0	1	0.38	0.00%	118.6/sec	0.00	0.00	0
S2 reached	100	0	0	1	0.38	0.00%	118.6/sec	0.00	0.00	0
TOTAL	200	3	0	22	4.25	0.00%	233.9/sec	42.51	24.90	186.1

## Task 2

1.

Sequence Diagram to show differences on stickiness We think that the SERVERID alternative is more effective because the balancer can just evaluate the SERVERID and redirect the request accordingly.

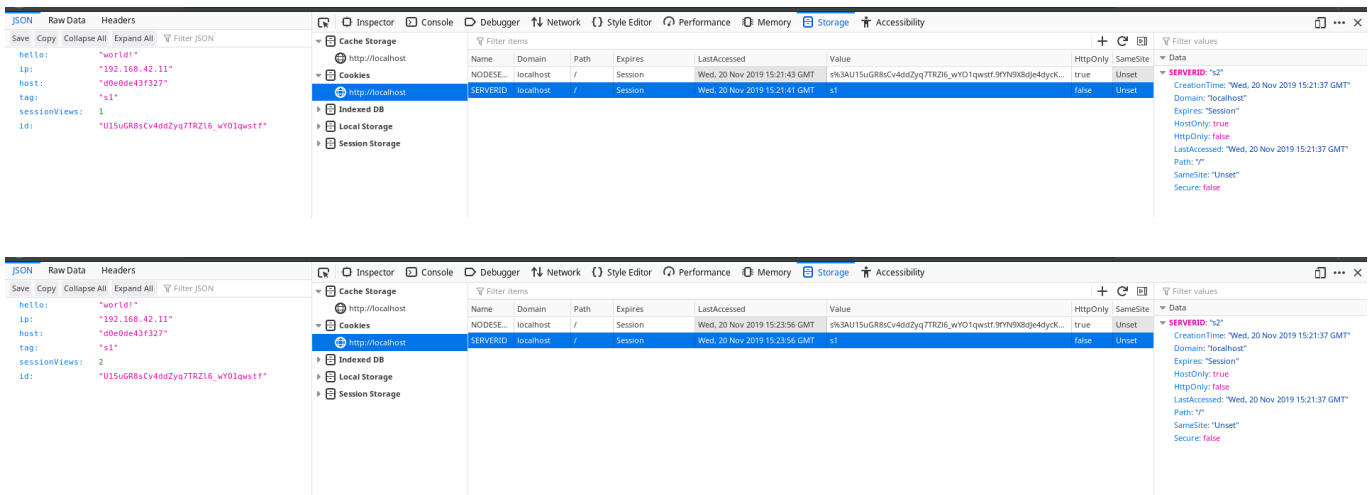
2.

We can see the modified configuration at the root of the git repo : `haproxy_sticky.cfg`. The modifications are juste on the `backend nodes` section. We've just added `cookie SERVERID insert indirect nocache` and the value we want to add on the cookie for each server so :

```
server s1 ${WEBAPP_1_IP}:3000 cookie s1 check
server s2 ${WEBAPP_2_IP}:3000 cookie s2 check
```

3.

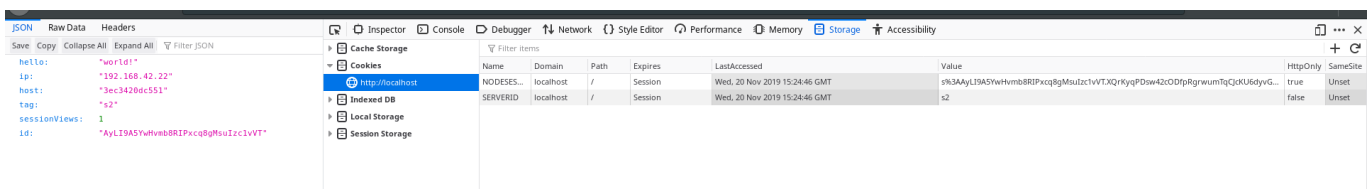
To see if we have achieved the right behaviour we can check like that : First we visit the site and we see that the load balancer crete the SERVERID cookie and that if we refresh the page we keep the same server :



The first screenshot shows the initial state of the browser's storage. The 'Cookies' section lists a single cookie named 'SERVERID' with the value 's1'. The 'Data' section shows the 'CreationTime' as 'Wed, 20 Nov 2019 15:21:37 GMT'.


The second screenshot shows the state after a page refresh. The 'Cookies' section still lists the 'SERVERID' cookie, but the value has changed to 's2'. The 'Data' section shows the 'CreationTime' as 'Wed, 20 Nov 2019 15:21:37 GMT'.

But if we open an incognito (to remove all cookies) we can see that the roundrobin strategy apply and we have the other server attributed :



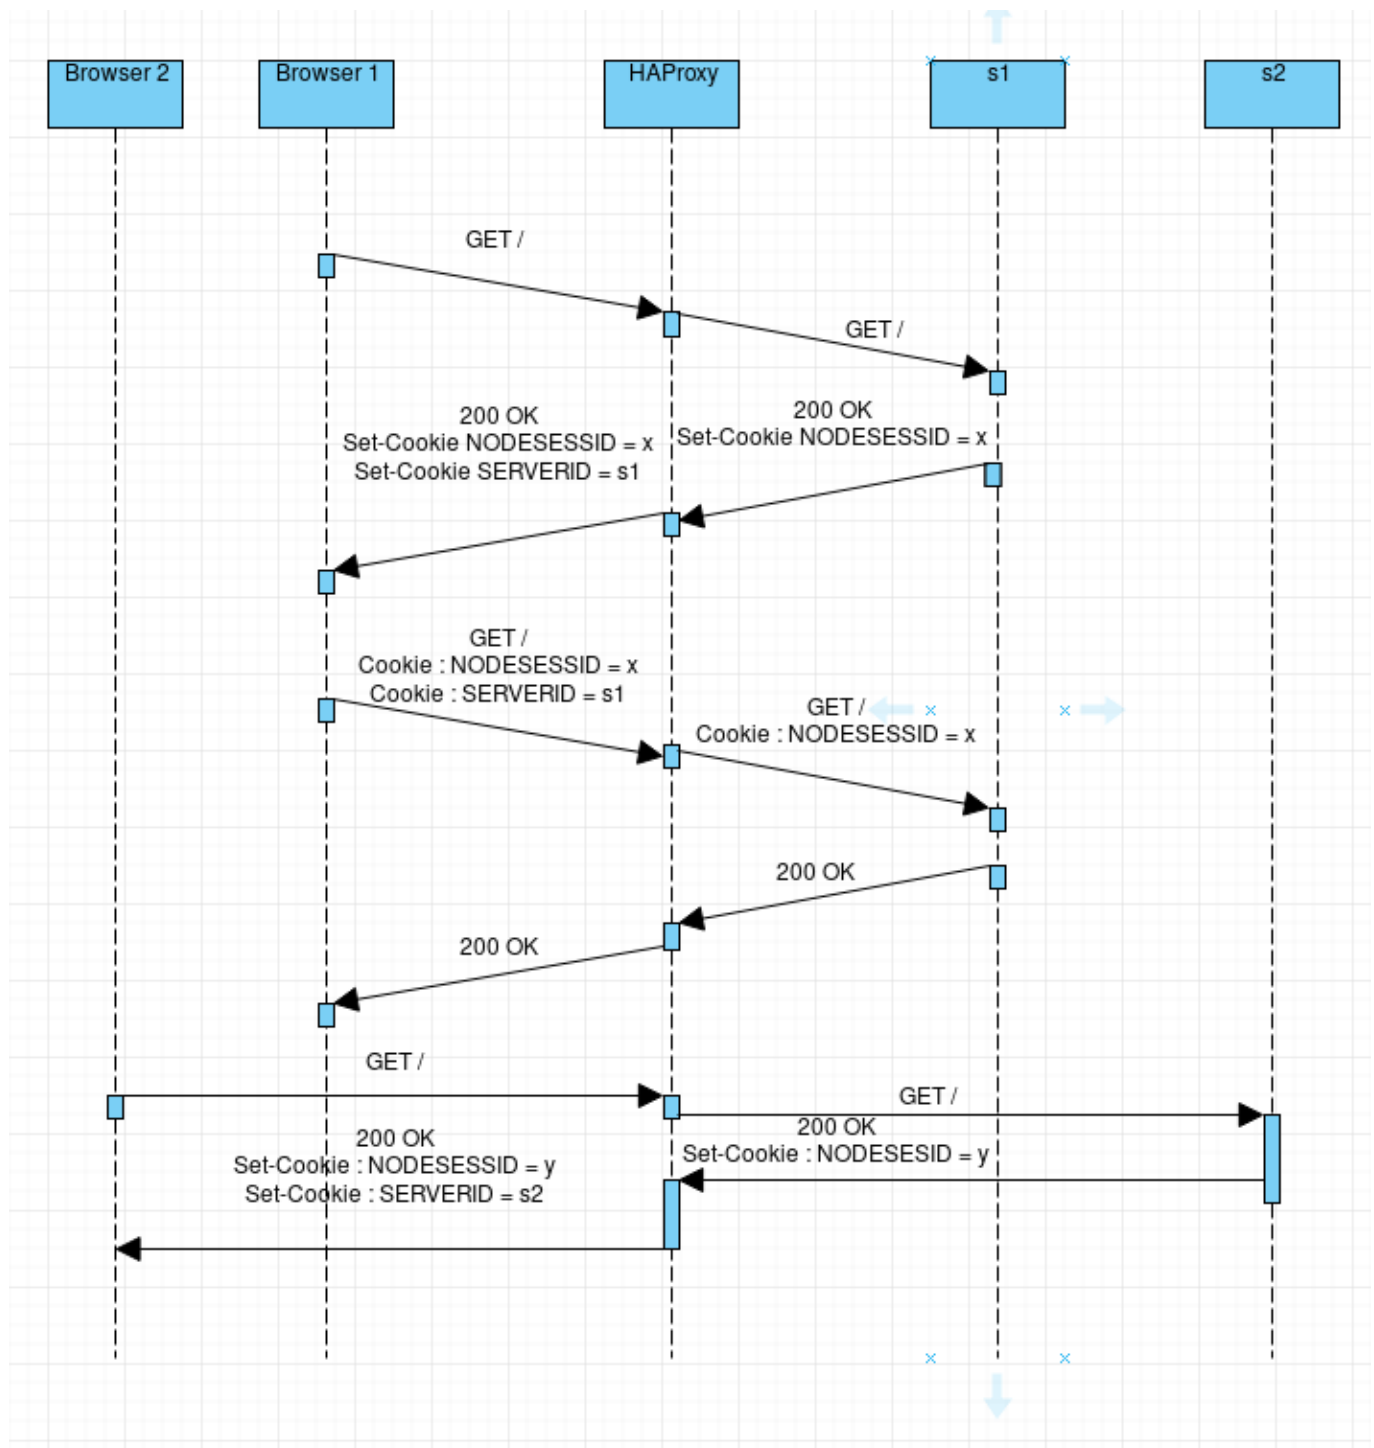
The screenshot shows the 'Cookies' section with the 'SERVERID' cookie having the value 's2'. The 'Data' section shows the 'CreationTime' as 'Wed, 20 Nov 2019 15:24:46 GMT'.

And if we refresh this page we keep the same server again because the SERVERID cookie is well set.

 sticky session 4

4.

We can see here the diagram of the situation where 1 browser refresh the page and what happens if another browser opens a connection :



We can see that the SERVERID cookie is set at the HAProxy level, and it's removed at this level too. And because of the roundrobin if another browser connects to the infra it will be directed to s2 if s1 was the last to respond. This cookie is used by the HAProxy to redirect correctly to the right server.

5.

Here is the summary report of this task :

Summary Report

Name: Summary Report

Comments:

Write results to file / Read from file

Filename

Browse...

Log/Display Only: ☐ Errors ☒ Successes 

Configure

Label	# Samples	Average	Min	Max	Std. Dev.	Error %	Throughput	Received KB/sec	Sent KB/sec	Avg. Bytes
GET /	1000	9	1	60	16.38	0.00%	81.7/sec	29.70	18.67	372.1
\$1 reached	1000	0	0	1	0.38	0.00%	81.8/sec	0.00	0.00	.0
TOTAL	2000	5	0	60	12.56	0.00%	163.5/sec	29.70	18.67	186.0

We can explain the behaviour because how JMeter handles the cookies, because there's only one thread group and so one user. So that's normal that we just have one server who responds (because of the SERVERID Cookie).

7.

We can see the result of the manipulations we have done, in fact we added one "user". So the second user when doing his request he's redirected to the other server and all of these future requests will too (because of the cookie).

Task 3

1.

HAProxy

Statistics Report for pid 10

> General process information

pid = 10 (process #1, nproc = 1)  
uptime = 50.0h34m01s  
systemd Units: nproc=1 unlimited; ulimit=1 = 4096  
maxrss=1 = 4096; maxconn = 2000; maxpipes = 0  
current ports = 1; current pipes = 0; conn rate = 0/sec  
Running tasks: 1/0 (0% = 0%)

active UP  
active UP going down  
active DOWN going up  
active or backup DOWN  
active or backup DOWN for maintenance  
active or backup SOFT STOPPED for maintenance  
Note: "NOLEAF/DRAIN" = UP with load balancing disabled

Display option:  
• Scope  
• Hide DOWN servers  
• Refresh now  
• CSV export

External resources:  
• Primary site  
• Upstream (L3/L4)  
• Online manual

Stats		Session rate		Sessions		Bytes		Errors		Warnings		Status		Server													
Queue	Cur	Max	Limit	Cur	Max	Limit	Total	Left	Last	In	Out	Req	Resp	Req	Conn	Resp	Rate	Redis	Status	LastChk	Uptime	Act	Back	Chk	Down	Downtime	Throttle
Frontend	0	1	-	1	2	2000	16			272 143	7 577 397	0	0	0	0	0	0	0	0	OPEN		0	0	0	0	0	
Backend	0	0	0	0	0	200	0	0	0s	272 143	7 577 397	0	0	0	0	0	0	0	0	34m1s UP		0	0	0	0	0	

Local nodes		Session rate		Sessions		Bytes		Errors		Warnings		Status		Server													
Queue	Cur	Max	Limit	Cur	Max	Limit	Total	Left	Last	In	Out	Req	Resp	Req	Conn	Resp	Rate	Redis	Status	LastChk	Uptime	Act	Back	Chk	Down	Downtime	Throttle
Frontend	0	1	-	1	2	2000	0			62 161	40 152	0	0	0	0	0	0	0	0	OPEN		0	0	0	0	0	

Nodes		Session rate		Sessions		Bytes		Errors		Warnings		Status		Server													
Queue	Cur	Max	Limit	Cur	Max	Limit	Total	Left	Last	In	Out	Req	Resp	Req	Conn	Resp	Rate	Redis	Status	LastChk	Uptime	Act	Back	Chk	Down	Downtime	Throttle
s1	0	0	-	0	0	0	1	-	3s	2	12m27s	10 057	11 305	0	0	0	0	0	0	33m57s UP	L7OK/200 in 4ms	1	Y	-	1	1	4s
s2	0	0	-	0	0	0	1	-	71	3	11m27s	47 104	28 840	0	0	0	0	0	0	34m1s UP	L7OK/200 in 2ms	1	Y	-	0	0	0s
Backend	0	0	0	0	0	1	200	99	0	11m27s	62 161	40 152	0	0	0	0	0	0	0	34m1s UP		2	2	0	0	0	0s

JSON

Données brutes

En-têtes

Enregistrer

Copier

Tout réduire

Tout développer

Filtrer le JSON

hello:

"world!"

ip:

"192.168.42.22"

host:

"66c5c8308b0e"

tag:

"s2"

sessionViews:

6

id:

"UxPOe0g9--4GcZt3RdiQnE6JmVkatYa"

We can see in the screenshot above that it is the node "s2" that answers.

2.

5 / 9

## HAProxy

### Statistics Report for pid 10

### > General process information

[illegible]

We pass the node "s2" on DRAIN mode, it become blue "soft stopped" on the state page.

3.

It's the same node than answer our request. In DRAIN mode, all the new traffic will be redirected to the other nodes, but the current session continue to make request to the node in DRAIN mode.

4.

```
{ "hello": "world!", "ip": "192.168.42.11", "host": "d2e8afe714f3", "tag": "s1", "sessionViews": 8, "id": "q9KMi56T2unOk1lK45rSdKZWd186a07Y" }
```

In the new browser , we start a new session and we reach the node "s1".

5.

```
{"hello":"world!","ip":"192.168.42.11","host":"d2e8afe714f3","tag":"s1","sessionViews":1,"id":"L1rfRh7i2LJ4bSqXVbOxQaIrguOvFFs"}
```

We clear the cookies on the new browser and we still reach the node "s1". Because "s2" is DRAIN mode, the new session can't reach the "s2". The new browser can only reach "s1".

```
JSON  Données brutes  En-têtes

Enregistrer Copier Tout réduire Tout développer Filtre le JSON

hello: "world!"
ip: "192.168.42.22"
host: "66c5c8308b0e"
tag: "s2"
sessionViews: 29
id: "UxPOe0g9--4GcZt3RdiCQNE6JmVkatYA"
```



6.

### Statistics Report for pid 10

[illegible]

It balanced sequentially between the two nodes.

```
{ "hello": "world!", "ip": "192.168.42.22", "host": "66c5c8308b0e", "tag": "s2", "sessionViews": 4, "id": "UPtIM ZOo3iSkHT7lso96WwMAn0AvYM0" }
```

The screenshot shows the JSON editor interface with the following components:

- Top Bar:** Contains tabs for "JSON" (selected), "Données brutes", and "En-têtes". Below these are action buttons: "Enregistrer", "Copier", "Tout réduire", "Tout développer", and a search icon with the text "Filtrer le JSON".
- JSON Content:** A JSON object is displayed with the following fields and values:
  - `hello:` `"world!"`
  - `ip:` `"192.168.42.22"`
  - `host:` `"66c5c8308b0e"`
  - `tag:` `"s2"`
  - `sessionViews:` `33`
  - `id:` `"UxPOe0g9--4GczT3RdiCQNE6JmvykatYa"`

7.

### Statistics Report for pid 10

General process information

pid = 10 (process #1, nbproc = 1)  
uptime = 0d 11h16m1s  
system limits: maxmem = unlimited, ulimits = 4094  
maxproc = 4094, maxopen = 2000, maxpipes = 0  
current open = 1, current pipes = 0.0, open rate = 1/sec  
Running tasks: 1/9 = 100 %

active UP

active UP, going down

active DOWN, going up

active or backup DOWN

active or backup DOWN for maintenance (MAINT)

active or backup STOP STOPPED for maintenance

backup UP

backup UP, going down

backup DOWN, going up

not checked

not checked

Note: "HOLD"/"DRAIN" = UP with load-balancing disabled

Display option:

Scope

Host DOWN status

Backup status

CRUI status

External resources

[PROMETHEUS](#)

[JOURNALS/LOGS](#)

[Online manual](#)

stats

	Queue			Session rate			Sessions			Bytes			Errors			Warnings			Status			Server									
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	Left	Last	In	Out	Req	Resp	Req	Conn	Resp	Ref	Redis	Status	LastCdn	Uptime	Act	Back	Chk	Down	Downtime	Throttle	
Frontend				1	2	-	1	2	-	2 000	72			290 558	0	2 220 977	0	0	0	0	0	0	116mUP								
Backend	0	0	0	0	0	0	0	0	0	300	0	0s		290 558	0	2 220 977	0	0	0	0	0	0									

localnodes

	Queue			Session rate			Sessions			Bytes			Errors			Warnings			Status			Server								
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	Left	Last	In	Out	Req	Resp	Req	Conn	Resp	Ref	Redis	Status	LastCdn	Uptime	Act	Back	Chk	Down	Downtime	Throttle
Frontend				0	1	-	0	3	2 000	103				99 058	87 411	0	0	0	0	0	0	OPEN								

nodes

	Queue			Session rate			Sessions			Bytes			Errors			Warnings			Status			Server								
	Cur	Max	Limit	Cur	Max	Limit	Cur	Max	Limit	Total	Left	Last	In	Out	Req	Resp	Req	Conn	Resp	Ref	Redis	Status	LastCdn	Uptime	Act	Back	Chk	Down	Downtime	Throttle
s1	0	0	0	0	0	1	-	63	8	3m24s	35 039		20 202		0	0	0	0	0	0	0	116mUP	L7OK/200 in 1mg	1	Y	-	1	1	4s	-
s2	0	0	0	0	0	1	-	102	4	3m12s	64 518		41 209		0	0	0	0	0	0	0	43sMAINT		1	Y	-	1	1	12mth	-

On MAINT mode, all the traffic is redirected to the other nodes, including existing sessions. The node "s2" can't be reached by request.

```
{"hello":"world!","ip":"192.168.42.11","host":"d2e8afe714f3","tag":"s1","sessionViews":1,"id":"yEsiGQ_Uz6Wgf-waB6wiGd7T6u_FmHlWw"}
```

The new browser reach the node "s1". If we clear the cookies, we still only reach "s1".

JSON	Données brutes	En-têtes
Enregistrer	Copier	Tout réduire
	Tout développer	Filtrer le JSON
hello:	"world!"	
ip:	"192.168.42.11"	
host:	"d2e8afe714f3"	
tag:	"s1"	
sessionViews:	1	
id:	"RWmemHV5130leKU_EvrES4-QRsIeerps"	

The first browser is not anymore on the node "s2". And has a new session on the node "s1".

## Task 4

1.

We have been asked to do a run for base data with the base config, so we set the delays to 0 for both servers (we took the JMeter conf. used for Task 2, 2 thread groups):

Label	# Samples	Average	Min	Max	Std. Dev.	Error %	Throughput	Received KB/sec	Sent KB/sec	Avg. Bytes
GET /	2000	16	2	58	19.36	0.00%	102.0/sec	37.08	23.41	372.1
S2 reached	1000	0	0	1	0.37	0.00%	52.0/sec	0.00	0.00	.0
S1 reached	1000	0	0	2	0.36	0.00%	52.4/sec	0.00	0.00	.0
TOTAL	4000	8	0	58	15.90	0.00%	204.1/sec	37.08	23.41	186.0

2.

We set a 250ms delay to s1 :

Label	# Samples	Average	Min	Max	Std. Dev.	Error %	Throughput	Received KB/sec	Sent KB/sec	Avg. Bytes
GET /	2000	763	1	2038	828.15	0.00%	1.3/sec	0.48	0.30	372.1
S2 reached	1000	0	0	4	0.51	0.00%	47.6/sec	0.00	0.00	.0
S1 reached	1000	0	0	6	0.50	0.00%	39.6/min	0.00	0.00	.0
TOTAL	4000	382	0	2038	699.02	0.00%	2.6/sec	0.48	0.30	186.0

We can see with this run taht it took a really long time (25min) because  $1000 \times 0.25s$ , but it works well in fact. What happens it's that the first user takes the connection to s1 and it will be long because the cookie specify s1 for each request. And the second user go to s2 without problems and takes each time s2 because of the cookie.

3.

Proof that we set correctly the delay :

POST /delay	2.276 ms - 47									
HEAD /	1.408 ms - 129									
HEAD /	1.500 ms - 129									
HEAD /	2506.530 ms - 129									
HEAD /	1.318 ms - 129									
HEAD /	1.525 ms - 129									
HEAD /	2508.473 ms - 129									
HEAD /	1.621 ms - 129									
HEAD /	1.685 ms - 129									
[WARNING] 335/123344 (11) : Server nodes/s1 is DOWN, reason: Layer7_timeout, check duration: 2001ms. 1 active and 0 backup servers left. 0 sessions active, 0 requested, 0 remaini										
HEAD /	2508.326 ms - 129									
HEAD /	1.722 ms - 129									
HEAD /	1.542 ms - 129									
HEAD /	2508.321 ms - 129									
HEAD /	1.358 ms - 129									

We can see on the proof that HA detects s1 as DOWN. We set a 2500ms delay to s1 :



Label	# Samples	Average	Min	Max	Std. Dev.	Error %	Throughput	Received KB/sec	Sent KB/sec	Avg. Bytes
GET /	2000	15	2	58	19.09	0.00%	103.9/sec	37.76	23.94	372.1
S2 reached	2000	0	0	2	0.38	0.00%	104.0/sec	0.00	0.00	.0
TOTAL	4000	8	0	58	15.63	0.00%	207.8/sec	37.75	23.93	186.0

We see that apparently the load balancer managed to see that the s1 server was too slow and directed all requests to s2 (see proof above). But that's really weird because this timeout for connecting to a server is defined as 5000ms.

4.

No we didn't have any error on these tasks, we think that the load balancer is smart enough to balance quite well, and we didn't make enough requests to cause an error.

5.

After doing the weight config we have this behaviour in the JMeter tests :

Label	# Samples	Average	Min	Max	Std. Dev.	Error %	Throughput	Received KB/sec	Sent KB/sec	Avg. Bytes
GET /	2000	155	1	570	149.98	0.00%	6.7/sec	2.42	1.52	372.1
S1 reached	1000	0	0	1	0.36	0.00%	3.3/sec	0.00	0.00	.0
S2 reached	1000	0	0	1	0.36	0.00%	57.8/sec	0.00	0.00	.0
TOTAL	4000	77	0	570	131.54	0.00%	13.3/sec	2.42	1.52	186.0

Nothing change from the 2. apart of the time taken. That's because the weights doesn't change anything for 2 users.

6.

Now the weights takes effects because it's like we have another user each time, because we clean cookies for each iteration. So the weights influence the behaviour of the load balancer in this way :

Label	# Samples	Average	Min	Max	Std. Dev.	Error %	Throughput	Received KB/sec	Sent KB/sec	Avg. Bytes
GET /	2000	373	3	1032	291.74	0.00%	5.2/sec	2.81	0.60	552.6
S1 reached	1334	0	0	1	0.39	0.00%	3.5/sec	0.00	0.00	.0
S2 reached	666	0	0	1	0.37	0.00%	1.7/sec	0.00	0.00	.0
TOTAL	4000	186	0	1032	278.17	0.00%	10.4/sec	2.81	0.60	276.3